

Big Data Infrastructures

Fall 2018

Lab 03 : HDFS + Hbase

Author : Thomas Schaller, Sylvain Julmy

Professor : Philippe Cudré-Mauroux

Task 1

We used the following command to copy the Weblogs to HDFS :

```
./bundle/hadoop/bin/hdfs dfs -copyFromLocal  
  ./weblogs_hbase.txt  
  /bdi_2018/bdi18_07/
```

Task 2

We used the following Java code to read the data from the Weblogs file and put it into the HBase Table.

```

14 public class CreateHBaseTable {
15
16     static Configuration conf = HBaseConfiguration.create();
17
18     public static void main(String[] args) throws Exception {
19
20         String[] months = {"Jan", "Feb", "Mar", "Apr", "May", "Jun", "Jul", "Aug",
21             ↪ "Sep", "Oct", "Nov", "Dec"};
22
23         conf.set("hbase.zookeeper.quorum", "diufm210.unifr.ch");
24         conf.set("hbase.zookeeper.property.clientPort", "2181");
25
26         Connection connection = ConnectionFactory.createConnection(conf);
27
28         String tableName = "weblogs_bdi18_07";
29         // Connect to Hbase and create the table
30         try {
31             Admin hAdmin = connection.getAdmin();
32             HTableDescriptor hTableDesc = new HTableDescriptor(
33                 TableName.valueOf(tableName));
34             hTableDesc.addFamily(new HColumnDescriptor("Months"));
35             hTableDesc.addFamily(new HColumnDescriptor("Statistics"));
36             System.out.println("Creating Table...");
37             hAdmin.createTable(hTableDesc);
38             System.out.println("Table created Successfully...");
39
40         } catch (Exception e) {
41             e.printStackTrace();
42         }
43
44         System.out.println("Connection to the table");
45         Table table = connection.getTable(TableName.valueOf(tableName));
46
47         try {
48
49             System.out.println("Insert the datas");
50             // Reading the dataset from HDFS
51             Path path = new Path("/bdi_2018/bdi18_07/weblogs_hbase.txt");
52             FileSystem fileSystem = FileSystem.get(new Configuration());
53             BufferedReader bufferedReader = new BufferedReader(new
54             ↪ InputStreamReader(fileSystem.open(path)));
55
56             String line = bufferedReader.readLine();
57             while (line != null) {

```

```

        String[] elements = line.split("\t");

```

```

58         assert elements.length == 13;
59
60         String rowKey = elements[0];
61
62         Put put = new Put(Bytes.toBytes(rowKey));
63
64         int totalNumberOfVisit = 0;
65
66         for (int i = 1; i < elements.length; i++) {
67             String key = months[i - 1];
68             String valueStr = elements[i];
69             int value = Integer.parseInt(valueStr);
70             if (value == 0)
71                 continue;
72             totalNumberOfVisit += value;
73             put.addColumn(
74                 Bytes.toBytes("Months"),
75                 Bytes.toBytes(key),
76                 Bytes.toBytes(valueStr)
77             );
78         }
79
80         put.addColumn(
81             Bytes.toBytes("Statistics"),
82             Bytes.toBytes("Active"),
83             Bytes.toBytes(
84                 totalNumberOfVisit > 20 ?
85                     "1" :
86                     "0"
87             )
88         );
89
90         table.put(put);
91
92         line = bufferedReader.readLine();
93     }
94
95     System.out.println("Done");
96     table.close();
97
98     } catch (IOException e) {
99         e.printStackTrace();
100     }
101
102 }
103

```

Task 3

The following sections contains the output of the specified command.

1)

Command : scan 'weblogs_bdi18_07'

Output :

```
hbase(main):014:0> scan 'weblogs_bdi18_07'
```

```
ROW COLUMN+CELL
 0.308.86.81|2012 column=Months:Jul, timestamp=1543329698909, value=1
 0.308.86.81|2012 column=Statistics:Active, timestamp=1543329698909, value=0
 0.32.48.676|2012 column=Months:Jan, timestamp=1543329698919, value=3
...
88.88.38.655|2012 column=Statistics:Active, timestamp=1543329752426, value=0
88.88.600.360|2012 column=Months:Jan, timestamp=1543329752427, value=2
88.88.600.360|2012 column=Statistics:Active, timestamp=1543329752427, value=0
88.88.618.331|2012 column=Months:May, timestamp=1543329752429, value=1
88.88.618.331|2012 column=Statistics:Active, timestamp=1543329752429, value=0
88.88.644.648|2012 column=Months:Jan, timestamp=1543329752431, value=1
88.88.644.648|2012 column=Statistics:Active, timestamp=1543329752431, value=0
88.88.67.600|2012 column=Months:Jul, timestamp=1543329752432, value=3
88.88.67.600|2012 column=Statistics:Active, timestamp=1543329752432, value=0
88.88.687.376|2012 column=Months:Jan, timestamp=1543329752434, value=2
88.88.687.376|2012 column=Statistics:Active, timestamp=1543329752434, value=0
27300 row(s) in 7.2460 seconds
```

2)

Command :

```
scan 'weblogs_bdi18_07', {  
  STARTROW => '0.32.85.668|2012',  
  ENDROW => '01.660.68.623|2012'  
}
```

Output :

```
hbase(main):023:0> scan 'weblogs_bdi18_07', {STARTROW => '0.32.85.668|2012', ENDROW => '01.660.68.623|2012'}  
ROW                                COLUMN+CELL  
  0.32.85.668|2012                column=Months:Jul, timestamp=1543329698922, value=8  
  0.32.85.668|2012                column=Statistics:Active, timestamp=1543329698922, value=0  
  0.45.305.7|2012                  column=Months:Feb, timestamp=1543329698926, value=1  
  0.45.305.7|2012                  column=Months:Jan, timestamp=1543329698926, value=1  
  0.45.305.7|2012                  column=Statistics:Active, timestamp=1543329698926, value=0  
  0.46.386.626|2011                column=Months:Nov, timestamp=1543329698928, value=1  
  0.46.386.626|2011                column=Statistics:Active, timestamp=1543329698928, value=0  
  0.48.322.75|2012                 column=Months:Jul, timestamp=1543329698931, value=1  
  0.48.322.75|2012                 column=Statistics:Active, timestamp=1543329698931, value=0  
  0.638.50.46|2011                 column=Months:Dec, timestamp=1543329698934, value=8  
  0.638.50.46|2011                 column=Statistics:Active, timestamp=1543329698934, value=0  
  0.87.36.333|2012                 column=Months:Aug, timestamp=1543329698937, value=7  
  0.87.36.333|2012                 column=Statistics:Active, timestamp=1543329698937, value=0  
6 row(s) in 0.0090 seconds
```

3)

Command : count 'weblogs_bdi18_07'

Output :

```
hbase(main):004:0> count 'weblogs_bdi18_07'
Current count: 1000, row: 11.638.80.681|2012
Current count: 2000, row: 14.676.84.33|2012
Current count: 3000, row: 18.614.66.380|2012
Current count: 4000, row: 322.05.67.601|2012
Current count: 5000, row: 323.55.374.668|2011
Current count: 6000, row: 325.83.602.85|2011
Current count: 7000, row: 328.327.620.3|2012
Current count: 8000, row: 362.4.40.321|2012
Current count: 9000, row: 366.387.680.320|2012
Current count: 10000, row: 41.388.661.660|2011
Current count: 11000, row: 44.81.54.615|2011
Current count: 12000, row: 48.681.648.08|2011
Current count: 13000, row: 52.682.638.604|2011
Current count: 14000, row: 55.30.58.687|2012
Current count: 15000, row: 57.68.658.31|2011
Current count: 16000, row: 606.41.1.88|2012
Current count: 17000, row: 630.630.322.65|2012
Current count: 18000, row: 638.38.386.658|2012
Current count: 19000, row: 651.05.680.613|2012
Current count: 20000, row: 658.624.85.64|2012
Current count: 21000, row: 668.302.304.308|2012
Current count: 22000, row: 680.686.17.85|2012
Current count: 23000, row: 682.674.56.58|2012
Current count: 24000, row: 687.624.84.684|2011
Current count: 25000, row: 80.331.62.07|2012
Current count: 26000, row: 85.610.688.8|2011
Current count: 27000, row: 88.630.610.80|2012
27300 row(s) in 2.9290 seconds
```

=> 27300

Task 4

Each following section present the code which perform the specified task. Each code is also available inside the zip archive.

1)

Task : Retrieve only the contents of the Columns: “Jan” and “Feb” from the row key: 06.305.307.336|2012.

Java code :

```
1 // Retrieve only the contents of the Columns:
2 // Jan and Feb from the row key: 06.305.307.336|2012
3
4 Get get1 = new Get(Bytes.toBytes("06.305.307.336|2012"));
5 Result result1 = table.get(get1);
6
7 String janValue = Bytes.toString(result1.getValue(
8     Bytes.toBytes("Months"),
9     Bytes.toBytes("Jan")
10 ));
11
12 String febValue = Bytes.toString(result1.getValue(
13     Bytes.toBytes("Months"),
14     Bytes.toBytes("Feb")
15 ));
16
17 System.out.printf("Jan : %s, Feb : %s \n", janValue, febValue);
18
19 System.out.println("Part 1 done");
```

2)

Task : Create a new ip and year, and fill in the table with the same values as the row with key: 01.660.70.74|2012

Java code :

```
1 //Create a new ip and year, and fill in the table with the same values as the row with key:
2 ↪ 01.660.70.74|2012
3
4 String newIp = "8.8.8.8";
5 String newYear = "1971";
6 String rowKey = newIp + "|" + newYear;
7
8 Put put = new Put(Bytes.toBytes(rowKey));
9
10 Get get2 = new Get(Bytes.toBytes("01.660.70.74|2012"));
11 Result result2 = table.get(get2);
```

```

12 NavigableMap<byte[], NavigableMap<byte[], NavigableMap<Long, byte[]>>> map = result2.getMap();
13
14 for (Map.Entry<byte[], NavigableMap<byte[], NavigableMap<Long, byte[]>>> familyEntry :
15     ↪ map.entrySet()) {
16
17     byte[] family = familyEntry.getKey();
18
19     for (Map.Entry<byte[], NavigableMap<Long, byte[]>> colEntry : map.get(family).entrySet()) {
20
21         byte[] qualifier = colEntry.getKey();
22
23         for (Map.Entry<Long, byte[]> timestampedValue :
24             ↪ map.get(family).get(qualifier).entrySet()) {
25
26             byte[] value = timestampedValue.getValue();
27
28             put.addColumn(family, qualifier, value);
29         }
30     }
31 }
32
33 table.put(put);
34
35 System.out.println("Part 2 done");

```

Here is a snapshot which proves that the new rowKey has been inserted with the correct values.

```

hbase(main):008:0> get 'weblogs_bdi18_07','01.660.70.74|2012'
COLUMN                                CELL
  Months:Jul                          timestamp=1543329698943, value=1
  Statistics:Active                    timestamp=1543329698943, value=0
2 row(s) in 0.0100 seconds

```

```

hbase(main):009:0> get 'weblogs_bdi18_07','8.8.8.8|1971'
COLUMN                                CELL
  Months:Jul                          timestamp=1543333405817, value=1
  Statistics:Active                    timestamp=1543333405817, value=0
2 row(s) in 0.0300 seconds

```

3)

Task : Delete the row with key: 88.88.324.601|2012

Java code :

```

1 //Delete the row with key: 88.88.324.601|2012
2
3 Delete delete = new Delete(Bytes.toBytes("88.88.324.601|2012"));

```

```
4  
5 table.delete(delete);  
6  
7 System.out.println("Part 3 done");
```

Here is the a snapshot of the two same query. We just run the code above in between them.

```
hbase(main):011:0> get 'weblogs_bdi18_07','88.88.324.601|2012'  
COLUMN                                CELL  
Months:Feb                            timestamp=1543329752411, value=20  
Months:Jan                            timestamp=1543329752411, value=37  
Months:Sep                            timestamp=1543329752411, value=74  
Statistics:Active                     timestamp=1543329752411, value=1  
4 row(s) in 0.0120 seconds
```

```
hbase(main):012:0> get 'weblogs_bdi18_07','88.88.324.601|2012'  
COLUMN                                CELL  
0 row(s) in 0.0060 seconds
```